

# SPECTRA: Faster Large Language Model Inference with Optimized Internal and External Speculation

Anonymous ACL submission

## Abstract

Inference with modern Large Language Models (LLMs) is both computationally expensive and time-consuming. Speculative decoding has emerged as a promising solution, but existing approaches face key limitations: training-based methods require a draft model that is challenging to obtain and lacks generalizability, while non-training methods offer limited speedup gains. In this work, we present SPECTRA, a novel framework for accelerating LLM inference without the need for additional training. SPECTRA introduces two innovative techniques for efficiently managing internal and external knowledge, each outperforming corresponding state-of-the-art (SOTA) methods independently. When combined, these techniques achieve up to a 4.08x speedup across various benchmarks and LLM architectures, significantly surpassing existing non-training approaches. The implementation of SPECTRA is publicly available.

## 1 Introduction

Generating long sequences with low latency using Large Language Models (LLMs) is a critical requirement. Current LLMs rely on autoregressive decoding (Touvron et al., 2023; Bai et al., 2023; Jiang et al., 2023; OpenAI et al., 2024), which suffers from inefficiency because it generates text one token at a time. This results in generation time scaling linearly with the sequence length and underutilizes the parallel processing capabilities of modern GPUs. A widely studied approach to mitigate this issue is speculative decoding (Chen et al., 2023; Leviathan et al., 2023), which follows a *guess-and-verify* paradigm. In this approach, a smaller LLM (draft model) (Chen et al., 2023; Leviathan et al., 2023; Miao et al., 2024; Sun et al., 2023b; Zhou et al., 2024; Cai et al., 2024) or the original LLM trained in a specialized manner (self-speculative decoding) (Elhoushi et al., 2024; Liu et al., 2024a; Yang et al., 2024; Zhang et al., 2024a;

Li et al., 2024b) predicts multiple tokens in advance. The original LLM then verifies these predictions in parallel, improving efficiency. However, these approaches require additional training, which demands substantial computational resources and may degrade the original model’s capabilities.

Another line of research focuses on speculating subsequent tokens without requiring additional training. This approach eliminates the need for training new models or modifying the original large language model (LLM), making it practical for off-the-shelf deployment. Some methods leverage specialized mechanisms to generate speculative tokens directly from the LLM’s predictions (Fu et al., 2024; Ou et al., 2024), while others rely on external information sources to derive these tokens (Yang et al., 2023; He et al., 2024; Li et al., 2024a). However, the speedup gain in these approaches remains limited due to the quality of the speculative guesses.

We introduce SPECTRA (Figure 1a), a speculative decoding method that improves generation speed without requiring any training or modifications to the original LLM. SPECTRA consists of two main components: a core module (SPECTRA-CORE, Figure 1c), which integrates seamlessly into LLMs in a plug-and-play manner, and an optional retrieval module (SPECTRA-RETRIEVAL, Figure 1e) that further enhances performance. The core module SPECTRA-CORE improves speculative decoding by leveraging the token distribution predicted by the LLM to generate high-quality guesses. Specifically, it employs two multi-level N-gram dictionaries that enable bi-directional search for dynamic-length guesses, balancing both quality and quantity. Additionally, SPECTRA optimizes a candidate pool to continuously update the N-gram dictionaries, ensuring broad token coverage. All updates to these resources, along with guess verification, are performed efficiently in a single forward pass. The retrieval module, SPECTRA-RETRIEVAL,

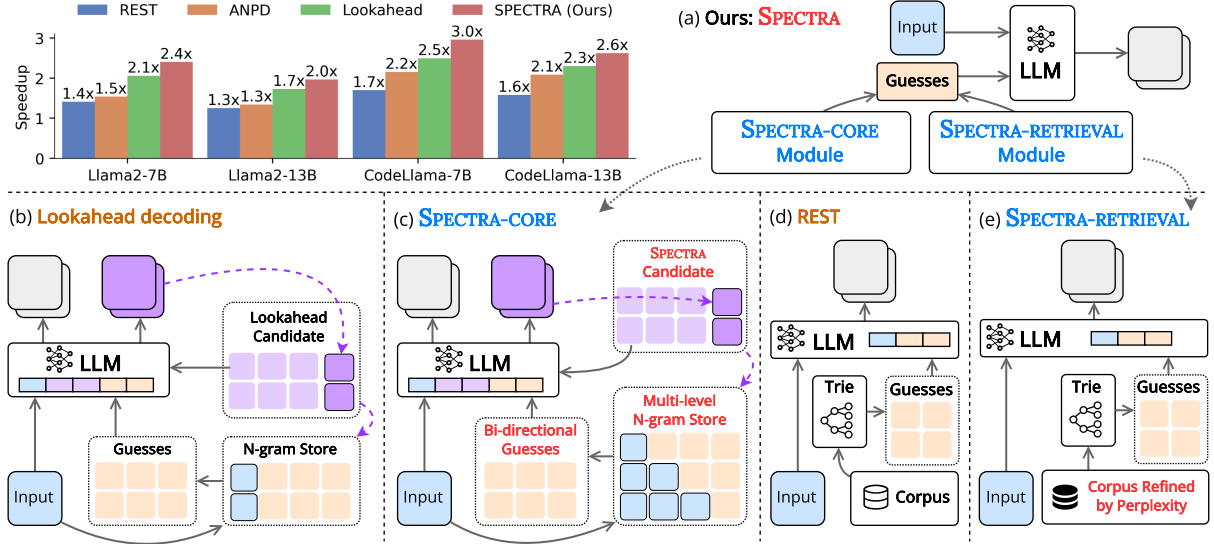


Figure 1: Overview of Spectra and comparison with other non-training SOTA approaches. (a) Overview of SPECTRA. (b) Overview of Lookahead Decoding (Fu et al., 2024). (c) Overview of the SPECTRA-CORE module, which utilizes the knowledge inside LLM for obtaining guesses. (d) Overview of REST (He et al., 2024). (e) Overview of the SPECTRA-RETRIEVAL module, which is designed to be integrated efficiently with SPECTRA-CORE to boost the speedup. The results in the bar chart are measured on HumanEval.

can be integrated to further enhance speedup. Existing approaches that rely on external sources for generating guesses (He et al., 2024) struggle to integrate with other speculative decoding methods, as the search time outweighs the speedup gains. SPECTRA-RETRIEVAL addresses this issue by reducing the search space, selecting only high-quality content from the corpus based on perplexity scores computed by the target LLM. This optimization enables seamless integration with SPECTRA-CORE, maximizing efficiency.

Empirical results on six tasks—including multi-turn conversation, code generation, and mathematical reasoning—across three LLM families (Llama 2 (Touvron et al., 2023), Llama 3 (Dubey et al., 2024), and CodeLlama (Rozière et al., 2024)) with model sizes ranging from 7B to 70B demonstrate that SPECTRA outperforms other non-training speculative decoding methods, achieving speedups of up to 4x. We publicly release the code and data. The key contributions of this paper are as follows:

- We introduce SPECTRA, which improves speculative decoding by effectively leveraging the LLM’s predicted token distribution. SPECTRA is a plug-and-play solution that requires no modifications to the LLM (Section 3.1).
- SPECTRA’s retrieval module refines external corpora using perplexity scores computed by the target LLM, providing a general frame-

work that enables speculative decoding approaches relying on external information to be seamlessly integrated with other speculative decoding techniques (Section 3.2).

- Extensive experiments across diverse tasks, LLM architectures, GPU types, and settings demonstrate the efficiency of SPECTRA, outperforming other non-training speculative decoding approaches (Section 5). SPECTRA also integrates with acceleration tools such as FlashAttention and pipeline parallelism (Section 5.2). The code and data are available.

## 2 Preliminaries

### 2.1 Autoregressive Decoding in LLMs

Given an input sequence  $\mathbf{x} = (x_1, x_2, \dots, x_s)$  of length  $s$ , and a slice of length  $m$  as  $\mathbf{x}_{1:m} = (x_1, x_2, \dots, x_m)$ , the output of an LLM represents a probability distribution over the next token. The probability of generating the  $s$ -th token, conditioned on all preceding tokens, is given by  $P_M(x_s | x_{1:s-1})$ . The next token  $x_s$  is sampled from this distribution using methods such as greedy, top- $k$ , or top- $p$  sampling (see (Kool et al., 2020; Holtzman et al., 2020)). For greedy sampling, the next token is selected as  $x_s = \arg \max P_M(x_s | x_{1:s-1})$ . Consequently, the LLM generates an output sequence  $(y_1, y_2, \dots, y_m)$  of length  $m$  autoregressively, where each token  $y_i$  is computed as

$$y_i = \operatorname{argmax} P_M(y_i \mid y_{1:i-1}, \mathbf{x}).$$

## 2.2 Speculative Decoding

Speculative decoding follows a *guess-and-verify* approach, where multiple candidate future tokens are speculated and subsequently verified in a single decoding step. With tree attention (Miao et al., 2024), multiple drafts can be verified simultaneously. Let  $G$  denote the number of guesses, and define the set of guesses as  $\tilde{Y} = \{\tilde{y}^{(0)}, \tilde{y}^{(1)}, \dots, \tilde{y}^{(G)}\}$ , where each guess sequence has length  $K$ . The  $j$ -th token of the  $i$ -th guess is denoted as  $\tilde{y}_j^{(i)}$ .

In the case of speculative decoding with greedy sampling, given the prompt  $\mathbf{x}$ , a drafting method generates the draft sequences  $\tilde{Y}$ . Using these drafts, the LLM computes the true tokens  $(y'_1, y'_2, \dots, y'_K)$  in parallel. These tokens are then verified, and  $h$  is defined as the highest number of correctly guessed tokens across all guesses. Consequently,  $h + 1$  tokens are generated in a single forward step. Algorithm 2 outlines speculative decoding with greedy sampling, and additional details are provided in Appendix A.

## 3 SPECTRA DECODING

SPECTRA consists of two modules (SPECTRA-CORE and SPECTRA-RETRIEVAL) that can function independently or together. The core module (SPECTRA-CORE) improves speedup by leveraging the LLM’s predicted token distribution to generate high-quality guesses and integrates into LLMs in a plug-and-play manner. The retrieval module (SPECTRA-RETRIEVAL) derives guesses from a refined external information source and is designed to integrate with SPECTRA-CORE to further enhance performance.

### 3.1 SPECTRA-CORE

SPECTRA-CORE maintains an  $N$ -gram storage and a candidate pool. The candidate pool  $\mathcal{C}$  contains  $W$  sequences,  $\{c^{(0)}, c^{(1)}, \dots, c^{(W-1)}\}$ , with each sequence consisting of  $N$  tokens. Let  $c_j^{(i)}$  represent the  $j$ -th token in the  $i$ -th sequence. The  $N$ -gram storage includes two dictionaries: the forward dictionary  $\mathcal{S}_{\text{fwd}}$  and the backward dictionary  $\mathcal{S}_{\text{bwd}}$ . At each time step, guesses  $\mathcal{G}$  are obtained through a bidirectional search using  $\mathcal{S}_{\text{fwd}}$  and  $\mathcal{S}_{\text{bwd}}$ . A single forward pass to the LLM retrieves all necessary distributions, which are used to generate new candidate tokens for  $\mathcal{C}$  and verify the guesses  $\mathcal{G}$ .

### Algorithm 1 SPECTRA Internal Knowledge

---

**Require:** Sequence  $\mathbf{x} = (x_1, x_2, \dots, x_n)$ , model  $P_M$ , max  $N$ -gram size  $N$ , candidate pool size  $W$ , max guesses  $G$ , max number of new tokens  $m$ . Refine threshold  $\tau$

- 1: Initialize  $N$ -gram Forward-dictionary  $\mathcal{S}_{\text{fwd}} \leftarrow \emptyset$
- 2: Initialize  $N$ -gram Backward-dictionary  $\mathcal{S}_{\text{bwd}} \leftarrow \emptyset$
- 3: Random  $c_j^{(i)}, \forall j \in [0, N-1], \forall i \in [0..W-1]$
- 4:  $t \leftarrow n+1$
- 5: **while**  $t \leq m$  **do**
- 6:   {Obtain the guesses}
- 7:    $\mathcal{G} \leftarrow \mathcal{S}_{\text{fwd}}[\mathbf{x}_{t-1}]$
- 8:    $u = \emptyset$
- 9:   **for**  $j = 0$  to  $N-1$  **do**
- 10:     **for**  $k = N-1$  to  $1$  **do**
- 11:        $u_j \leftarrow \mathcal{S}_{\text{bwd}}[\mathbf{x}_{t+j-k:t-1} \oplus u_{0:j-1}]$
- 12:       break if found value for  $u_j$
- 13:     **end for**
- 14:   **end for**
- 15:    $\mathcal{G}.\text{append}(u)$
- 16:   {Forward in LLM}
- 17:   Obtain necessary distributions of  $P_M$  in parallel.
- 18:   {Verification}
- 19:   {Greedy verify (Alg. 3) or Sampling verify (Alg. 4)}
- 20:    $\text{hits} \leftarrow \text{VerificationFunction}(\mathbf{x}, P_M, \mathbf{g})$
- 21:    $\mathbf{x} = \mathbf{x} \oplus \text{hits}$
- 22:    $t \leftarrow t + \text{size}(\text{hits})$
- 23:   {Predict Candidates}
- 24:   **for**  $i = 0$  to  $W-1$  **do**
- 25:      $r \sim \text{Uniform}[0, 1]$
- 26:      $P_c(c_{N-1}^{(i)}) \leftarrow P_M(c_{N-1}^{(i)} \mid c_{:N-2}^{(i)}, \mathbf{x})$
- 27:     **if**  $r > \tau$  **then**
- 28:        $c_{N-1}^{(i)} \leftarrow \operatorname{argmax}_{c \notin \mathcal{S}_{\text{fwd}}} P_c(c_{N-1}^{(i)})$
- 29:     **else**
- 30:        $c_{N-1}^{(i)} \leftarrow \operatorname{argmax} P_c(c_{N-1}^{(i)})$
- 31:     **end if**
- 32:   **end for**
- 33:   {Update  $N$ -gram dictionaries}
- 34:   **for**  $i = 0$  to  $W-1$  **do**
- 35:     **for**  $j = 0$  to  $N-2$  **do**
- 36:        $\mathcal{S}_{\text{fwd}}[c_j^{(i)}].\text{append}(c_{j+1}^{(i)})$
- 37:        $\mathcal{S}_{\text{bwd}}[c_{0:j}^{(i)}] \leftarrow c_{j+1}^{(i)}$
- 38:     **end for**
- 39:   **end for**
- 40:   {Update Candidates}
- 41:    $c_j^{(i)} \leftarrow c_{j+1}^{(i)}, \forall j \in [0, N-2], \forall i$
- 42: **end while**
- 43: **Output:**  $\mathbf{x}_{n+1:n+m} = (y_1, y_2, \dots, y_m)$

---

The dictionaries  $\mathcal{S}_{\text{fwd}}$  and  $\mathcal{S}_{\text{bwd}}$  are updated with  $N$ -grams from the candidate pool. The details of the SPECTRA-CORE decoding process are described in Algorithm 1.

**Bi-directional Search for Guesses** At each step, SPECTRA generates  $G$  guess sequences  $\mathcal{G} = \{\tilde{y}^{(0)}, \tilde{y}^{(1)}, \dots, \tilde{y}^{(G)}\}$ . Unlike previous work (Fu et al., 2024), which enforces uniform guess lengths, SPECTRA supports variable-length guesses, improving both flexibility and efficiency. The forward dictionary  $\mathcal{S}_{\text{fwd}}$  maps a token to a *list of sequences*, while the backward dictionary  $\mathcal{S}_{\text{bwd}}$  maps a sequence to a single token. At time step




Figure 2: Details of SPECTRA forward step in LLM. The dashed arrow indicates interactions between the tokens, which are realized by the LLM’s attention mask.

$t$ , the set of guesses is obtained through a bidirectional search (Alg. 1, lines 7–15). This search operates in two directions: (1) the forward direction, which prioritizes the quantity of guesses, and (2) the backward direction, which prioritizes the quality of guesses. In the forward direction, the last generated token  $x_{t-1}$  is used to search  $\mathcal{S}_{fwd}$  for guess sequences (Alg. 1, line 7). In the backward direction, a high-quality guess is constructed by iteratively predicting one token at a time using  $\mathcal{S}_{bwd}$ , repeating the process until a desired sequence length  $N$  is reached (Alg. 1, lines 8–14).

**Predict & Verify in One Forward Pass** All distributions required for predicting candidates and verifying guesses are obtained in a single forward pass to the LLM, leveraging parallel processing (Figure 2). This is achieved using a specially designed attention mask that specifies the allowed interactions between tokens. For instance, the token  $c_2^{(1)}$  attends only to  $c_1^{(1)}$ ,  $c_0^{(1)}$ , and the input.

**Predict Tokens for Candidate Pool** We predict the next candidate tokens  $c_{N-1}^{(i)}$  for the candidate pool using the distribution obtained from the forward pass (Alg. 1, lines 24–32). A straightforward approach is to select tokens with the highest probability in the token distribution. However, we observe that when searching for guesses in the forward dictionary  $\mathcal{S}_{fwd}$ , it is crucial for the search token to exist in the dictionary; otherwise, no guesses can be retrieved. To address this, we introduce a randomness-based mechanism to increase the coverage of  $\mathcal{S}_{fwd}$ . Specifically, we probabilistically encourage the selection of unseen tokens in  $\mathcal{S}_{fwd}$  using a hyperparameter  $\tau \in [0, 1]$ . Let  $r$  be a random draw from  $[0, 1]$ . If  $r > \tau$ , we select tokens with the highest probability that are not in  $\mathcal{S}_{fwd}$ ; otherwise, we choose tokens with the highest probability regardless of their presence in  $\mathcal{S}_{fwd}$ . Although  $c_{N-1}^{(i)}$  does not immediately affect the

coverage of  $\mathcal{S}_{fwd}$ , it contributes to coverage expansion in subsequent time steps through our candidate updating mechanism. At the end of each time step, all candidate sequences are shifted left by one token:  $c_j^{(i)} \leftarrow c_{j+1}^{(i)}$ , leaving  $c_{N-1}^{(i)}$  empty and ready for prediction in the next time step (Alg. 1, line 41).

**Update N-gram Dictionaries** At the end of each time step, candidate tokens from the pool  $\mathcal{C}$  are used to update the N-gram dictionaries  $\mathcal{S}_{fwd}$  and  $\mathcal{S}_{bwd}$ . While previous work (Fu et al., 2024) only adds the full N-gram  $(c_0^{(i)}, c_1^{(i)}, \dots, c_N^{(i)})$ , we observe that subsequences within N-grams often appear later in the generation process. By including these subsequences in the N-gram storage, we improve both the quality of guesses and the coverage of the dictionaries. Specifically, we add subsequences to  $\mathcal{S}_{fwd}$  using the first token as the key, and update  $\mathcal{S}_{bwd}$  by mapping the preceding part of the sequence to the last token (Alg. 1, lines 33–39).

### 3.2 SPECTRA-RETRIEVAL

SPECTRA-RETRIEVAL leverages an external knowledge source to generate guesses. This involves processing a text corpus and indexing it into a structure that supports fast prefix search, such as a trie. At each time step, the last generated tokens are used as input to this structure to retrieve guesses for speculative decoding. However, we observe that using random texts from the corpus without selection can limit the speedup gain. To address this, we propose a method to identify and select high-quality, relevant texts from the corpus tailored to the specific LLM. This improves the speedup gain and enables seamless integration with other speculative decoding approaches, including SPECTRA-CORE.

**Corpus Refinement by Perplexity** Given a text sequence  $u = (u_0, u_1, u_2, \dots)$ , perplexity quantifies the average uncertainty of the model when predicting the next token, conditioned on the preceding tokens. It is calculated as:

$$\text{PPL}(u) = \exp \left\{ -\frac{1}{t} \sum_{i=1}^t \log p_{\theta}(u_i | u_{<i}) \right\}$$

A lower perplexity indicates that the model assigns higher probabilities to the sequence, suggesting that the sequence is well-aligned with the model’s predictions and can produce high-quality guesses for speculative decoding. To optimize the retrieval process, we select texts with the lowest perplexity



from the corpus to form a smaller, high-quality subset, which is then used to construct the trie structure for generating guesses.

**Integration with SPECTRA-CORE** Our experiments (Section 5.2, Table 2) demonstrate that naively integrating guesses from external sources (e.g., REST (He et al., 2024)) into other speculative methods (e.g., Lookahead (Fu et al., 2024)) can lead to a noticeable drop in speedup. This occurs because the forward pass in the LLM can only handle a limited number of guesses, and exceeding this limit increases memory usage and slows down generation. With a limited guess budget, guesses from external sources can only account for a fraction of the total guesses, causing the search time in the indexing structure (e.g., a trie) to outweigh the speedup gain. To address this, it is crucial to limit the size of the external knowledge while maintaining the quality of the guesses. By refining the corpus using perplexity, SPECTRA-RETRIEVAL seamlessly integrates with SPECTRA-CORE, further boosting the speedup gain. Specifically, we integrate SPECTRA-RETRIEVAL into SPECTRA-CORE by including its guesses in the set of guesses during the guess generation step (Alg. 1, lines 7–15).

## 4 Experiments

**Models.** We evaluate LLaMA-2-Chat 7B, 13B, 70B (Touvron et al., 2023), CodeLlama 7B, 13B (Rozière et al., 2024), and LLaMA-3-Instruct 8B, 70B (Dubey et al., 2024).

**Tasks.** We conduct comprehensive evaluations on various generation tasks. MT-Bench (Zheng et al., 2023) for multi-turn conversation; GSM8K (Cobbe et al., 2021) for mathematical reasoning; HumanEval (Chen et al., 2021), MBPP (Austin et al., 2021) and ClassEval (Du et al., 2023) for code generation.

**Metrics.** SPECTRA does not modify the original LLM and the acceptance conditions, making it a lossless acceleration method. Therefore, the generation quality remains the same as the original LLM. We only evaluate the acceleration performance using the following metrics.

- **Speedup Ratio:** The speedup ratio relative to autoregressive decoding.
- **Compression ratio:** The ratio of the total number of autoregressive steps to the number

of Spectra decoding steps needed to produce the same sequence length.

**Baselines.** We use standard autoregressive decoding as the baseline (speed-up ratio = 1.00x). We further compare SPECTRA with leading non-training speculative decoding approaches, namely Adaptive N-gram (Ou et al., 2024), REST (He et al., 2024), and Lookahead (Fu et al., 2024). For details regarding implementation settings of both SPECTRA and these baselines, please refer to Appendix B.

## 5 Results

### 5.1 Main Results

**Overall Performance.** The top portion of Table 1 presents the speedup ratios of all evaluated methods under a greedy decoding setup. Our approach, SPECTRA, consistently yields the highest acceleration across the entire range of datasets and LLMs. In particular, SPECTRA achieves speedups up to  $4.08\times$  with LLaMA-3-8B-Instruct on the MBPP dataset.

For smaller models (7B), SPECTRA often surpasses  $3\times$  acceleration, underscoring the effectiveness of multi-token compression. By contrast, for 13B models, while the boost remains strong, it is relatively more moderate, typically falling in the  $1.6\times$ – $3\times$  band. We attribute this trend to the increased overhead of each forward pass in larger networks, which can dampen the proportional gains of fewer decoding iterations per token. Despite this, SPECTRA continues to outperform baselines across all parameter settings.

Significant advantages are evident in tasks such as GSM8K and ClassEval, where outputs often follow recurring patterns (e.g., repeated variable names or class definitions). In these scenarios, SPECTRA combines internal knowledge of partial sequences with external retrieval suggestions, thereby proposing accurate multi-token guesses. On the other hand, in domains featuring more varied or unpredictable responses—such as complex multi-turn conversations in MT-Bench—the acceptance rate is somewhat lower, although still competitive.

**Compression Ratio.** Table 1 also reports each method’s compression rate, a measure agnostic to specific hardware configurations. Across every dataset and LLM tested, SPECTRA delivers the highest average compression ratio. Each of SPECTRA’s draft-and-verify iterations typically yields

Model	Method	Classeval		GSM8K		Humaneval		MBPP		MTBench		AVG
		speedup	$\tau$	speedup	$\tau$	speedup	$\tau$	speedup	$\tau$	speedup	$\tau$	speedup
Greedy (temperature=0)												
CL-13B	ANPD	1.94	2.52	2.81	3.72	2.08	2.50	2.71	3.58	2.61	3.41	2.43
CL-13B	Lookahead	2.25	3.61	2.80	4.24	2.30	3.16	2.91	4.44	2.59	4.04	2.57
CL-13B	REST	1.28	2.14	0.93	1.54	1.58	2.31	0.85	1.40	0.94	1.53	1.12
CL-13B	SPECTRA (Ours)	<b>2.38</b>	<b>4.06</b>	<b>2.91</b>	<b>4.65</b>	<b>2.63</b>	<b>3.95</b>	<b>3.29</b>	<b>4.46</b>	<b>2.65</b>	<b>4.40</b>	<b>2.77</b>
CL-7B	ANPD	2.30	2.68	3.21	3.75	2.16	2.47	3.16	3.78	3.35	3.83	2.84
CL-7B	Lookahead	2.59	3.66	2.99	3.83	2.50	3.05	2.90	3.67	3.23	4.27	2.84
CL-7B	REST	1.45	2.22	0.91	1.39	1.70	2.34	0.96	1.45	1.02	1.44	1.21
CL-7B	SPECTRA (Ours)	<b>2.70</b>	<b>4.10</b>	<b>3.33</b>	<b>4.59</b>	<b>2.96</b>	<b>3.90</b>	<b>3.56</b>	<b>4.45</b>	<b>3.70</b>	<b>4.52</b>	<b>3.25</b>
L2-13B	ANPD	1.36	1.78	1.47	1.72	1.34	1.61	1.12	1.32	1.17	1.37	1.29
L2-13B	Lookahead	1.81	2.76	1.46	1.87	1.73	2.32	1.38	1.69	1.51	2.04	1.58
L2-13B	REST	1.22	2.01	0.94	1.46	1.25	1.94	0.95	1.44	1.14	1.90	1.10
L2-13B	SPECTRA (Ours)	<b>2.00</b>	<b>3.24</b>	<b>1.83</b>	<b>2.62</b>	<b>1.96</b>	<b>2.91</b>	<b>1.63</b>	<b>2.24</b>	<b>1.75</b>	<b>2.60</b>	<b>1.83</b>
L2-70B	ANPD	1.82	1.90	1.63	1.61	1.86	1.87	1.17	1.20	1.34	1.30	1.56
L2-70B	Lookahead	2.65	2.87	1.86	2.02	2.57	2.67	1.49	1.54	1.94	2.00	2.10
L2-70B	SPECTRA (Ours)	<b>3.10</b>	<b>3.40</b>	<b>2.52</b>	<b>2.69</b>	<b>3.22</b>	<b>3.37</b>	<b>1.86</b>	<b>1.93</b>	<b>2.43</b>	<b>2.51</b>	<b>2.62</b>
L2-7B	ANPD	1.62	1.95	1.52	1.68	1.54	1.67	1.19	1.33	1.30	1.37	1.43
L2-7B	Lookahead	2.19	2.94	1.66	1.93	2.06	2.42	1.46	1.69	1.73	2.05	1.82
L2-7B	REST	1.36	2.12	1.01	1.47	1.41	2.04	1.01	1.46	1.25	1.90	1.21
L2-7B	SPECTRA (Ours)	<b>2.40</b>	<b>3.43</b>	<b>2.11</b>	<b>2.64</b>	<b>2.40</b>	<b>3.05</b>	<b>1.77</b>	<b>2.16</b>	<b>2.02</b>	<b>2.59</b>	<b>2.14</b>
L3-70B	ANPD	1.54	1.67	1.50	1.47	1.83	1.88	1.46	1.41	1.23	1.23	1.51
L3-70B	Lookahead	2.40	2.62	1.54	1.58	2.56	2.70	1.43	1.45	1.76	1.86	1.94
L3-70B	SPECTRA (Ours)	<b>2.67</b>	<b>2.91</b>	<b>2.10</b>	<b>2.14</b>	<b>2.84</b>	<b>3.02</b>	<b>1.94</b>	<b>1.94</b>	<b>2.06</b>	<b>2.13</b>	<b>2.32</b>
L3-8B	ANPD	2.11	2.49	3.86	4.57	1.83	2.09	3.36	3.58	1.14	1.23	2.46
L3-8B	Lookahead	2.59	3.44	3.71	4.61	2.49	2.89	3.79	4.65	1.53	1.85	2.82
L3-8B	SPECTRA (Ours)	<b>2.83</b>	<b>3.49</b>	<b>3.89</b>	<b>4.77</b>	<b>2.57</b>	<b>3.02</b>	<b>4.08</b>	<b>4.76</b>	<b>1.69</b>	<b>2.10</b>	<b>3.01</b>
Sampling (temperature=1.0)												
CL-13B	ANPD	1.15	1.46	1.07	1.31	1.05	1.30	1.00	1.24	2.31	2.89	1.31
CL-13B	Lookahead	1.38	2.00	1.08	1.43	1.29	1.75	1.02	1.34	2.33	3.48	1.42
CL-13B	REST	1.14	1.87	0.82	1.35	1.27	1.96	0.84	1.39	0.93	1.50	1.00
CL-13B	SPECTRA (Ours)	<b>1.68</b>	<b>2.22</b>	<b>1.20</b>	<b>1.75</b>	<b>1.65</b>	<b>2.12</b>	<b>1.15</b>	<b>1.70</b>	<b>2.37</b>	<b>3.80</b>	<b>1.61</b>
CL-7B	ANPD	1.29	1.50	1.16	1.30	1.10	1.32	1.12	1.27	2.77	3.05	1.49
CL-7B	Lookahead	1.54	2.03	1.19	1.41	1.43	1.81	1.19	1.43	2.72	3.50	1.61
CL-7B	REST	1.23	1.86	0.88	1.33	1.33	1.98	0.91	1.40	0.97	1.44	1.06
CL-7B	SPECTRA (Ours)	<b>1.81</b>	<b>2.25</b>	<b>1.35</b>	<b>1.73</b>	<b>1.68</b>	<b>2.12</b>	<b>1.33</b>	<b>1.72</b>	<b>2.78</b>	<b>3.94</b>	<b>1.79</b>
L2-13B	ANPD	1.20	1.52	1.24	1.46	1.17	1.40	1.03	1.22	1.17	1.35	1.16
L2-13B	Lookahead	1.52	2.22	1.32	1.69	1.48	2.00	1.18	1.48	1.49	2.01	1.40
L2-13B	REST	1.18	1.96	0.93	1.45	1.19	1.88	0.92	1.44	1.12	1.88	1.07
L2-13B	SPECTRA (Ours)	<b>1.70</b>	<b>2.75</b>	<b>1.55</b>	<b>2.23</b>	<b>1.69</b>	<b>2.59</b>	<b>1.34</b>	<b>1.89</b>	<b>1.74</b>	<b>2.57</b>	<b>1.60</b>
L2-7B	ANPD	1.31	1.51	1.34	1.48	1.28	1.46	1.10	1.22	1.25	1.36	1.26
L2-7B	Lookahead	1.78	2.30	1.51	1.76	1.72	2.09	1.25	1.49	1.68	2.02	1.59
L2-7B	REST	1.26	2.03	0.99	1.46	1.27	1.93	0.96	1.41	1.21	1.88	1.14
L2-7B	SPECTRA (Ours)	<b>1.97</b>	<b>2.83</b>	<b>1.78</b>	<b>2.28</b>	<b>2.04</b>	<b>2.75</b>	<b>1.47</b>	<b>1.84</b>	<b>1.97</b>	<b>2.54</b>	<b>1.85</b>
L3-8B	ANPD	1.25	1.37	1.97	2.18	1.43	1.65	1.89	2.07	1.15	1.21	1.54
L3-8B	Lookahead	1.48	1.78	2.07	2.41	1.79	2.21	1.99	2.40	1.57	1.81	1.78
L3-8B	SPECTRA (Ours)	<b>1.94</b>	<b>2.84</b>	<b>2.27</b>	<b>2.78</b>	<b>1.92</b>	<b>2.51</b>	<b>2.19</b>	<b>2.78</b>	<b>1.70</b>	<b>2.05</b>	<b>2.01</b>

Table 1: Overall performance of speculative decoding methods across multiple tasks. “CL- $x$ B” denotes CodeLlama with  $x$ B parameters, “L2- $x$ B” denotes LLaMA-2-Chat of size  $x$ B, and “L3- $x$ B” denotes LLaMA-3-Instruct of size  $x$ B. We report the speedup ratio (vs. autoregressive) and the compression ratio  $\tau$ .

2.1–4.8 tokens, substantially outpacing alternative approaches and nearly doubling the acceptance length achieved by REST.

**Acceleration in Sampling Decoding.** The lower section of Table 1 investigates the performance of SPECTRA under sampling-based decoding with a temperature of 1.0. The results highlight how

SPECTRA continues to accelerate generation relative to baselines, offering roughly  $1.15\text{--}2.77\times$  speedups over standard autoregressive decoding. These gains are more modest than in greedy decoding, reflecting the lower acceptance rate under the sampling-based verification phase, which is consistent with earlier findings (Fu et al., 2024; Leviathan et al., 2023).

## 5.2 Analysis

**Ablation study.** We conducted a detailed component-wise analysis to determine the contribution of each module to the framework’s overall performance (Table 2). Specifically, the results on LLaMA2-7B-chat reveal that removing different components yields varying impacts on GSM8K speedups. Under the “CORE Module” configuration, excluding multi-level  $n$ -grams lowers the speedup from  $2.04\times$  to  $1.95\times$  (a 4% decrease), whereas turning off forward information reduces it from  $2.04\times$  to  $1.50\times$  (a 26% drop). Similarly, omitting backward information results in a speedup of  $1.94\times$ , down from  $2.04\times$ . In contrast, the “RETRIEVAL Module” setting shows that leaving out perplexity-based filtering decreases the speedup from  $1.18\times$  to  $1.16\times$ . Our fully integrated approach, SPECTRA, achieves a  $2.14\times$  speedup on GSM8K—outperforming both the “CORE Module” ( $2.04\times$ ) and “RETRIEVAL Module” ( $1.18\times$ ) variants. This improvement demonstrates the importance of combining multi-level  $n$ -grams, forward/backward drafting, and perplexity-based refinement in boosting acceptance rates and enhancing overall speedups. A similar trend was also observed in the results of the MTBench dataset.

Additionally, we compared our method against a naive combination of Lookahead and REST—where guess sequences from REST are added to Lookahead. This combined approach falls significantly short of our SPECTRA method, highlighting that a simple merger of two techniques is insufficient without our carefully optimized integration strategy and components.

**Priority for source of guesses** Since verifying too many candidate tokens at once can strain GPU resources and reduce speedups (Fu et al., 2024; Li et al., 2024b), SPECTRA limits how many guesses proceed to verification in each step (Appendix B). To understand whether internal or external guesses are more valuable, we temporarily remove this cap and measure acceptance rates (Figure 3). We

Method	GSM8K		MTBench	
	speedup	$\tau$	speedup	$\tau$
REST	1.01	1.47	1.25	1.90
Lookahead	1.66	1.93	1.73	2.05
Lookahead + REST	1.08	1.47	1.27	1.90
SPECTRA’s ablation				
CORE Module	2.04	2.50	1.92	2.35
- w/o forward info	1.50	1.68	1.20	1.37
- w/o backward info	1.94	2.21	1.74	2.12
- w/o Sub-Ngram	1.95	2.34	1.75	2.18
RETRIEVAL Module	1.18	1.31	1.24	1.50
- w/o PPL refine	1.16	1.29	1.20	1.45
SPECTRA (ours)	<b>2.14</b>	<b>2.64</b>	<b>2.02</b>	<b>2.59</b>

Table 2: Ablation of SPECTRA’s components (greedy decoding, LLaMA2-7B-Chat). “Sub-Ngram” augments each  $n$ -gram with its sub-sequences; “forward/backward info” uses internal expansions; and “PPL refine” applies perplexity-based filtering for external retrieval. “Lookahead + REST” denotes a naive combination where guess sequences from REST are directly added to Lookahead

observe that sequences generated via internal expansions—particularly forward and backward predictions—have a higher acceptance probability than those retrieved from external sources. Consequently, SPECTRA prioritizes internal guesses for verification. Interestingly, in code-generation tasks like HumanEval, external suggestions become more influential, likely due to code’s repetitive structure and the retrieval of similar snippets. This observation indicates that a strategic blend of backward internal knowledge and external retrieval can be particularly fruitful in these domains, especially when computational resources limit extensive forward expansions.

**FlashAttention.** Figure 3 shows that enabling FlashAttention consistently boosts the speedup of all methods, albeit to varying degrees. Notably, we observe an additional  $0.24\times$  speedup gain for SPECTRA on both GSM8K and MTBench. This is because FlashAttention better exploits the parallel structure of speculative decoding by reducing attention overheads, especially when verifying multiple guessed tokens in parallel. Although smaller gains are also seen for other methods, SPECTRA benefits the most, as it presents the longest verification branches and thus stands to profit significantly from more efficient attention implementations.

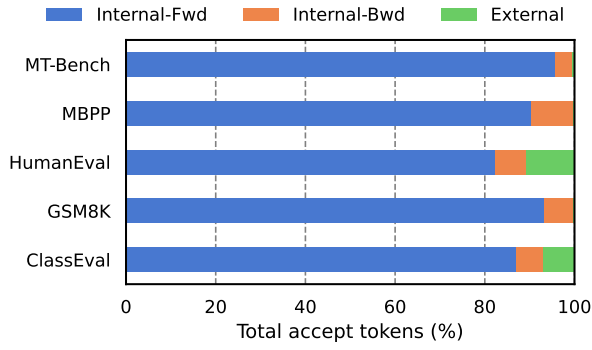


Figure 3: Acceptance rates for different guess sources (e.g., SPECTRA-CORE forward dictionary, backward dictionary, SPECTRA-RETRIEVAL’s guesses). The acceptance rate is the fraction of guessed tokens that pass verification and are appended to the final output.




Figure 4: Effect of FlashAttention on speculative decoding speed: Measured speedups on GSM8K and MTBench (LLama2-7B-Chat, greedy decoding). “No Flash” uses standard attention; “With Flash” uses FlashAttention for faster parallel verification.

while maintaining the exact output distribution (Chen et al., 2023; Leviathan et al., 2023). Some speculative decoding approaches train a smaller LLM (referred to as a draft model) (Chen et al., 2023; Leviathan et al., 2023; Miao et al., 2024; Sun et al., 2023b; Zhou et al., 2024; Cai et al., 2024), or train the original LLM itself in a special manner (referred to as self-speculative) (Elhoushi et al., 2024; Liu et al., 2024a; Yang et al., 2024; Zhang et al., 2024a; Li et al., 2024b) to guess several subsequent tokens and then verify them parallelly using the original LLM. As these approaches require training, they pose limitations, such as requiring heavy computational resources and losing the original model capabilities.

To avoid additional training, alternative speculative decoding methods leverage external resources or structural properties of language generation. Retrieval-based methods sidestep draft model training by using a datastore indexed with observed prefixes to retrieve guess sequences (Yang et al., 2023; He et al., 2024; Li et al., 2024a). Other approaches, such as Jacobi-like parallel decoding (Santilli et al., 2023) and lookahead decoding (Fu et al., 2024), mitigate left-to-right dependencies by generating and validating multiple candidate tokens in parallel. These training-free techniques achieve comparable speedups to learned methods without requiring model optimization, making them ideal for scenarios with computational or deployment constraints.

## 6 Related Works

Large language models (LLMs) are increasingly deployed in a range of applications, motivating ongoing research into more efficient inference (Liu et al., 2025). Common strategies include quantizing model weights into lower-precision formats (Liu et al., 2024b; Lin et al., 2024; Zhao et al., 2024; Park et al., 2024), pruning redundant parameters (Ma et al., 2023; Xia et al., 2023; Sun et al., 2023a; Le et al., 2025), and employing knowledge distillation (Gu et al., 2024; Friha et al., 2024; Zhang et al., 2024b). These techniques help reduce the computational load per forward pass, thereby lowering generation latency. However, they often introduce some degradation in model performance, forcing practitioners to balance quality with efficiency.

A growing line of work explores *speculative decoding* as a strategy for accelerating generation

## 7 Conclusions

In this work, we introduced SPECTRA, a hybrid speculative decoding framework that combines multi-level  $n$ -grams (internal knowledge) with perplexity-based retrieval (external knowledge) to achieve speedups of up to  $4.08\times$  across various LLMs and benchmarks, without additional training or compromising exact output fidelity. Our ablation studies show that each module (multi-level  $n$ -grams, forward/backward expansions, and perplexity-based datastore curation) substantially boosts acceptance rates, and their synergy outperforms existing non-training methods. By offering a lossless speedup that efficiently exploits both internal patterns and external texts, SPECTRA provides a practical, high-impact solution for accelerating inference in large language models.



## 8 Limitations

**(1) Cost of Building External Datastores.** Although our internal-knowledge strategy only relies on sequences observed during generation (and thus requires no extra data), our external-knowledge approach depends on constructing and indexing a sizeable datastore from potentially large corpora. This process can be time-consuming and memory-intensive, particularly in domains where data updates frequently or storage is constrained. While this additional investment can yield substantial speedups by increasing token acceptance rates, it may not be universally feasible or cost-effective.

**(2) Limited Evaluation Scope.** Our experiments center primarily on English-language benchmarks in conversational and coding tasks using LLaMA-based models. Although SPECTRA can, in principle, be applied to other models or languages, additional factors such as domain-specific tokenization or specialized textual structures may affect the acceptance rate and overall speedup. Future work is needed to assess the generality of SPECTRA across diverse linguistic settings (e.g., low-resource languages or specialized technical documents) and for a wider range of model families (beyond LLaMA-based architectures) to confirm and refine its applicability.

## References

Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and others. 2021. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*.

Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. 2023. Qwen Technical Report.

Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D. Lee, Deming Chen, and Tri Dao. 2024. MEDUSA: Simple LLM inference acceleration framework with multiple decoding heads. In *Proceedings of the 41st International Conference on*

*Machine Learning*, ICML’24. JMLR.org. Place: Vienna, Austria.

Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. 2023. *Accelerating Large Language Model Decoding with Speculative Sampling*.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, and others. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, and others. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.

Ning Ding, Yulin Chen, Bokai Xu, Yujia Qin, Shengding Hu, Zhiyuan Liu, Maosong Sun, and Bowen Zhou. 2023. *Enhancing Chat Language Models by Scaling High-quality Instructional Conversations*. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 3029–3051, Singapore. Association for Computational Linguistics.

Xueying Du, Mingwei Liu, Kaixin Wang, Hanlin Wang, Junwei Liu, Yixuan Chen, Jiayi Feng, Chaofeng Sha, Xin Peng, and Yiling Lou. 2023. Classeval: A manually-crafted benchmark for evaluating llms on class-level code generation. *arXiv preprint arXiv:2308.01861*.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, Ahmed Aly, Beidi Chen, and Carole-Jean Wu. 2024. *LayerSkip: Enabling Early Exit Inference and Self-Speculative Decoding*. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12622–12642, Bangkok, Thailand. Association for Computational Linguistics.

Othmane Friha, Mohamed Amine Ferrag, Burak Kantarci, Burak Cakmak, Arda Ozgun, and Nassira Ghoualmi-Zine. 2024. Llm-based edge intelligence: A comprehensive survey on architectures, applications, security and trustworthiness. *IEEE Open Journal of the Communications Society*.

Yichao Fu, Peter Bailis, Ion Stoica, and Hao Zhang. 2024. Break the sequential dependency of LLM inference using LOOKAHEAD DECODING. In *Proceedings of the 41st International Conference on*

633	<i>Machine Learning</i> , ICML'24. JMLR.org. Place: Vienna, Austria.	689
634		690
635	Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. 2024.	691
636	Minillm: Knowledge distillation of large language	692
637	models. In <i>The Twelfth International Conference on</i>	693
638	<i>Learning Representations</i> .	694
639	Zhenyu He, Zexuan Zhong, Tianle Cai, Jason Lee, and	695
640	Di He. 2024. <a href="#">REST: Retrieval-Based Speculative</a>	696
641	<a href="#">Decoding</a> . In <i>Proceedings of the 2024 Conference of</i>	697
642	<i>the North American Chapter of the Association for</i>	
643	<i>Computational Linguistics: Human Language Tech-</i>	698
644	<i>nologies (Volume 1: Long Papers)</i> , pages 1582–1595,	699
645	Mexico City, Mexico. Association for Computational	700
646	Linguistics.	701
647	Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and	702
648	Yejin Choi. 2020. <a href="#">The Curious Case of Neural Text</a>	703
649	<a href="#">Degeneration</a> . In <i>International Conference on Learn-</i>	
650	<i>ing Representations</i> .	704
651	Albert Q. Jiang, Alexandre Sablayrolles, Arthur Men-	705
652	sch, Chris Bamford, Devendra Singh Chaplot, Diego	706
653	de las Casas, Florian Bressand, Gianna Lengyel, Guil-	707
654	laume Lample, Lucile Saulnier, L��lio Renard Lavaud,	708
655	Marie-Anne Lachaux, Pierre Stock, Teven Le Scao,	709
656	Thibaut Lavril, Thomas Wang, Timoth��e Lacroix,	710
657	and William El Sayed. 2023. Mistral 7B.	711
658	Denis Kocetkov, Raymond Li, Loubna Ben allal, Jia LI,	
659	Chenghao Mou, Yacine Jernite, Margaret Mitchell,	712
660	Carlos Mu��oz Ferrandis, Sean Hughes, Thomas Wolf,	713
661	Dzmitry Bahdanau, Leandro Von Werra, and Harm de	714
662	Vries. 2023. <a href="#">The Stack: 3 TB of permissively li-</a>	715
663	<a href="#">censed source code</a> . <i>Transactions on Machine Learn-</i>	716
664	<i>ing Research</i> .	717
665	Wouter Kool, Herke van Hoof, and Max Welling. 2020.	718
666	<a href="#">Ancestral Gumbel-Top-k Sampling for Sampling</a>	719
667	<a href="#">Without Replacement</a> . <i>Journal of Machine Learning</i>	
668	<i>Research</i> , 21(47):1–36.	720
669	Khang Nguyen Le, Ryo Sato, Dai Nakashima, Takeshi	721
670	Suzuki, and Minh Le Nguyen. 2025. Optiprune: Ef-	722
671	fective pruning approach for every target sparsity. In	723
672	<i>Proceedings of the 31st International Conference on</i>	
673	<i>Computational Linguistics</i> , pages 3600–3612.	724
674	Yaniv Leviathan, Matan Kalman, and Yossi Matias.	725
675	2023. Fast inference from transformers via spec-	726
676	ulative decoding. In <i>Proceedings of the 40th Interna-</i>	727
677	<i>tional Conference on Machine Learning</i> , ICML'23.	728
678	JMLR.org. Place: Honolulu, Hawaii, USA.	729
679	Minghan Li, Xilun Chen, Ari Holtzman, Beidi Chen,	730
680	Jimmy Lin, Wen-tau Yih, and Xi Victoria Lin. 2024a.	731
681	<a href="#">Nearest Neighbor Speculative Decoding for LLM</a>	732
682	<a href="#">Generation and Attribution</a> . In <i>The Thirty-eighth An-</i>	733
683	<i>nuual Conference on Neural Information Processing</i>	734
684	<i>Systems</i> .	735
685	Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang	
686	Zhang. 2024b. <a href="#">EAGLE-2: Faster Inference of Lan-</a>	736
687	<a href="#">guage Models with Dynamic Draft Trees</a> . In <i>Proceed-</i>	737
688	<i>ings of the 2024 Conference on Empirical Methods</i>	738
	<i>in Natural Language Processing</i> , pages 7421–7432,	739
	Miami, Florida, USA. Association for Computational	740
	Linguistics.	741
	Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-	742
	Ming Chen, Wei-Chen Wang, Guangxuan Xiao,	743
	Xingyu Dang, Chuang Gan, and Song Han. 2024.	744
	Awq: Activation-aware weight quantization for on-	
	device llm compression and acceleration. <i>Proceed-</i>	
	<i>ings of Machine Learning and Systems</i> , 6:87–100.	
	Fangcheng Liu, Yehui Tang, Zhenhua Liu, Yunsheng	
	Ni, Duyu Tang, Kai Han, and Yunhe Wang. 2024a.	
	<a href="#">Kangaroo: Lossless Self-Speculative Decoding for</a>	
	<a href="#">Accelerating LLMs via Double Early Exiting</a> . In <i>The</i>	
	<i>Thirty-eighth Annual Conference on Neural Informa-</i>	
	<i>tion Processing Systems</i> .	
	Yiheng Liu, Hao He, Tianle Han, Xu Zhang, Mengyuan	
	Liu, Jiaming Tian, Yutong Zhang, Jiaqi Wang, Xi-	
	aohui Gao, Tianyang Zhong, Yi Pan, Shaochen Xu,	
	Zihao Wu, Zhengliang Liu, Xin Zhang, Shu Zhang,	
	Xintao Hu, Tuo Zhang, Ning Qiang, Tianming Liu,	
	and Bao Ge. 2025. <a href="#">Understanding llms: A compre-</a>	
	<a href="#">hensive overview from training to inference</a> . <i>Neuro-</i>	
	<i>computing</i> , 620:129190.	
	Zechun Liu, Barlas Oguz, Changsheng Zhao, Ernie	
	Chang, Pierre Stock, Yashar Mehdad, Yangyang Shi,	
	Raghuraman Krishnamoorthi, and Vikas Chandra.	
	2024b. <a href="#">LLM-QAT: Data-free quantization aware</a>	
	<a href="#">training for large language models</a> . In <i>Findings of</i>	
	<i>the Association for Computational Linguistics: ACL</i>	
	<i>2024</i> , pages 467–484, Bangkok, Thailand. Associa-	
	tion for Computational Linguistics.	
	Xinyin Ma, Gongfan Fang, and Xinchao Wang. 2023.	
	Llm-pruner: On the structural pruning of large lan-	
	guage models. <i>Advances in neural information pro-</i>	
	<i>cessing systems</i> , 36:21702–21720.	
	Xupeng Miao, Gabriele Oliaro, Zhihao Zhang, Xinhao	
	Cheng, Zeyu Wang, Zhengxin Zhang, Rae Ying Yee	
	Wong, Alan Zhu, Lijie Yang, Xiaoxiang Shi, Chunan	
	Shi, Zhuoming Chen, Daiyaan Arfeen, Reyna Ab-	
	hyankar, and Zhihao Jia. 2024. <a href="#">SpecInfer: Accelerat-</a>	
	<a href="#">ing Large Language Model Serving with Tree-based</a>	
	<a href="#">Speculative Inference and Verification</a> . In <i>Proceed-</i>	
	<i>ings of the 29th ACM International Conference on Ar-</i>	
	<i>chitectural Support for Programming Languages and</i>	
	<i>Operating Systems, Volume 3, ASPLOS '24</i> , pages	
	932–949, New York, NY, USA. Association for Com-	
	puting Machinery. Event-place: La Jolla, CA, USA.	
	Ramesh Nallapati, Bowen Zhou, Cicero dos Santos,	
	��a��lar Gul��ehre, and Bing Xiang. 2016. <a href="#">Abstrac-</a>	
	<a href="#">tive text summarization using sequence-to-sequence</a>	
	<a href="#">RNNs and beyond</a> . In <i>Proceedings of the 20th</i>	
	<i>SIGNLL Conference on Computational Natural Lan-</i>	
	<i>guage Learning</i> , pages 280–290, Berlin, Germany.	
	Association for Computational Linguistics.	
	Shashi Narayan, Shay B Cohen, and Mirella Lap-	
	ata. 2018. Don't give me the details, just the	

745	summary! topic-aware convolutional neural net-	
746	works for extreme summarization. <i>arXiv preprint</i>	
747	<i>arXiv:1808.08745</i> .	
748	OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal,	
749	Lama Ahmad, Ilge Akkaya, Florencia Leoni Ale-	
750	man, Diogo Almeida, Janko Altmenschmidt, Sam Alt-	
751	man, Shyamal Anadkat, Red Avila, Igor Babuschkin,	
752	Suchir Balaji, Valerie Balcom, Paul Baltescu, Haim-	
753	ing Bao, Mohammad Bavarian, Jeff Belgum, Ir-	
754	wan Bello, Jake Berdine, Gabriel Bernadett-Shapiro,	
755	Christopher Berner, Lenny Bogdonoff, Oleg Boiko,	
756	Madelaine Boyd, Anna-Luisa Brakman, Greg Brock-	
757	man, Tim Brooks, Miles Brundage, Kevin Button,	
758	Trevor Cai, Rosie Campbell, Andrew Cann, Brittany	
759	Carey, Chelsea Carlson, Rory Carmichael, Brooke	
760	Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully	
761	Chen, Ruby Chen, Jason Chen, Mark Chen, Ben	
762	Chess, Chester Cho, Casey Chu, Hyung Won Chung,	
763	Dave Cummings, Jeremiah Currier, Yunxing Dai,	
764	Cory Decareaux, Thomas Degry, Noah Deutsch,	
765	Damien Deville, Arka Dhar, David Dohan, Steve	
766	Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti,	
767	Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix,	
768	Simón Posada Fishman, Juston Forte, Isabella Ful-	
769	ford, Leo Gao, Elie Georges, Christian Gibson, Vik	
770	Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-	
771	Lopes, Jonathan Gordon, Morgan Grafstein, Scott	
772	Gray, Ryan Greene, Joshua Gross, Shixiang Shane	
773	Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris,	
774	Yuchen He, Mike Heaton, Johannes Heidecke, Chris	
775	Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele,	
776	Brandon Houghton, Kenny Hsu, Shengli Hu, Xin	
777	Hu, Joost Huizinga, Shantanu Jain, Shawn Jain,	
778	Joanne Jang, Angela Jiang, Roger Jiang, Haozhun	
779	Jin, Denny Jin, Shino Jomoto, Billie Jonn, Hee-	
780	woo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Ka-	
781	mali, Ingmar Kanitscheider, Nitish Shirish Keskar,	
782	Tabarak Khan, Logan Kilpatrick, Jong Wook Kim,	
783	Christina Kim, Yongjik Kim, Jan Hendrik Kirch-	
784	ner, Jamie Kiros, Matt Knight, Daniel Kokotajlo,	
785	Łukasz Kondraciuk, Andrew Kondrich, Aris Kon-	
786	stantinidis, Kyle Kopic, Gretchen Krueger, Vishal	
787	Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan	
788	Leike, Jade Leung, Daniel Levy, Chak Ming Li,	
789	Rachel Lim, Molly Lin, Stephanie Lin, Mateusz	
790	Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue,	
791	Anna Makanju, Kim Malfacini, Sam Manning, Todor	
792	Markov, Yaniv Markovski, Bianca Martin, Katie	
793	Mayer, Andrew Mayne, Bob McGrew, Scott Mayer	
794	McKinney, Christine McLeavey, Paul McMillan,	
795	Jake McNeil, David Medina, Aalok Mehta, Jacob	
796	Menick, Luke Metz, Andrey Mishchenko, Pamela	
797	Mishkin, Vinnie Monaco, Evan Morikawa, Daniel	
798	Mossing, Tong Mu, Mira Murati, Oleg Murk, David	
799	Mély, Ashvin Nair, Reiichiro Nakano, Rajeiv Nayak,	
800	Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh,	
801	Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex	
802	Paino, Joe Palermo, Ashley Pantuliano, Giambat-	
803	tista Parascandolo, Joel Parish, Emy Parparita, Alex	
804	Passos, Mikhail Pavlov, Andrew Peng, Adam Perel-	
805	man, Filipe de Avila Belbute Peres, Michael Petrov,	
806	Henrique Ponde de Oliveira Pinto, Michael, Poko-	
807	rny, Michelle Pokrass, Vitchyr H. Pong, Tolly Pow-	
	ell, Alethea Power, Boris Power, Elizabeth Proehl,	808
	Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh,	809
	Cameron Raymond, Francis Real, Kendra Rimbach,	810
	Carl Ross, Bob Rotsted, Henri Roussez, Nick Ry-	811
	der, Mario Saltarelli, Ted Sanders, Shibani Santurkar,	812
	Girish Sastry, Heather Schmidt, David Schnurr, John	813
	Schulman, Daniel Selsam, Kyla Sheppard, Toki	814
	Sherbakov, Jessica Shieh, Sarah Shoker, Pranav	815
	Shyam, Szymon Sidor, Eric Sigler, Maddie Simens,	816
	Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin	817
	Sokolowsky, Yang Song, Natalie Staudacher, Fe-	818
	lipe Petroski Such, Natalie Summers, Ilya Sutskever,	819
	Jie Tang, Nikolas Tezak, Madeleine B. Thompson,	820
	Phil Tillet, Amin Tootoonchian, Elizabeth Tseng,	821
	Preston Tuggle, Nick Turley, Jerry Tworek, Juan Fe-	822
	lipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya,	823
	Chelsea Voss, Carroll Wainwright, Justin Jay Wang,	824
	Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei,	825
	C. J. Weinmann, Akila Welihinda, Peter Welin-	826
	der, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave	827
	Willner, Clemens Winter, Samuel Wolrich, Hannah	828
	Wong, Lauren Workman, Sherwin Wu, Jeff Wu,	829
	Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin	830
	Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers,	831
	Chong Zhang, Marvin Zhang, Shengjia Zhao, Tian-	832
	hao Zheng, Juntang Zhuang, William Zhuk, and Bar-	833
	ret Zoph. 2024. <a href="#">GPT-4 Technical Report</a> .	834
	Jie Ou, Yueming Chen, and Prof. Tian. 2024. <a href="#">Lossless</a>	835
	<a href="#">Acceleration of Large Language Model via Adap-</a>	836
	<a href="#">tive N-gram Parallel Decoding</a> . In <i>Proceedings of</i>	837
	<i>the 2024 Conference of the North American Chap-</i>	838
	<i>ter of the Association for Computational Linguistics:</i>	839
	<i>Human Language Technologies (Volume 6: Industry</i>	840
	<i>Track)</i> , pages 10–22, Mexico City, Mexico. Associa-	841
	tion for Computational Linguistics.	842
	Yeonhong Park, Jake Hyun, SangLyul Cho, Bonggeun	843
	Sim, and Jae W. Lee. 2024. Any-precision llm: Low-	844
	cost deployment of multiple, different-sized llms. In	845
	<i>Proceedings of the 41st International Conference on</i>	846
	<i>Machine Learning</i> .	847
	Baptiste Rozière, Jonas Gehring, Fabian Gloeckle, Sten	848
	Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi,	849
	Jingyu Liu, Romain Sauvestre, Tal Remez, Jérémy	850
	Rapin, Artyom Kozhevnikov, Ivan Evtimov, Joanna	851
	Bitton, Manish Bhatt, Cristian Canton Ferrer, Aaron	852
	Grattafiori, Wenhan Xiong, Alexandre Défossez,	853
	Jade Copet, Faisal Azhar, Hugo Touvron, Louis Mar-	854
	tin, Nicolas Usunier, Thomas Scialom, and Gabriel	855
	Synnaeve. 2024. <a href="#">Code Llama: Open Foundation</a>	856
	<a href="#">Models for Code</a> . <i>_eprint</i> : 2308.12950.	857
	Andrea Santilli, Silvio Severino, Emilian Postolache,	858
	Valentino Maiorca, Michele Mancusi, Riccardo	859
	Marin, and Emanuele Rodola. 2023. <a href="#">Accelerating</a>	860
	<a href="#">transformer inference for translation via parallel de-</a>	861
	<a href="#">coding</a> . In <i>Proceedings of the 61st Annual Meeting of</i>	862
	<i>the Association for Computational Linguistics (Vol-</i>	863
	<i>ume 1: Long Papers)</i> , pages 12336–12355, Toronto,	864
	Canada. Association for Computational Linguistics.	865
	Mingjie Sun, Zhuang Liu, Anna Bair, and J. Zico Kolter.	866



2023a. [A simple and effective pruning approach for large language models](#). *ArXiv*, abs/2306.11695.

Ziteng Sun, Ananda Theertha Suresh, Jae Hun Ro, Ahmad Beirami, Himanshu Jain, and Felix Yu. 2023b. SpecTr: fast speculative decoding via optimal transport. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS '23, Red Hook, NY, USA. Curran Associates Inc. Event-place: New Orleans, LA, USA.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. Llama 2: Open Foundation and Fine-Tuned Chat Models.

Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, and Danqi Chen. 2023. [Sheared llama: Accelerating language model pre-training via structured pruning](#). *ArXiv*, abs/2310.06694.

Nan Yang, Tao Ge, Liang Wang, Binxing Jiao, Daxin Jiang, Linjun Yang, Rangan Majumder, and Furu Wei. 2023. [Inference with Reference: Lossless Acceleration of Large Language Models](#). *eprint*: 2304.04487.

Seongjun Yang, Gibbeum Lee, Jaewoong Cho, Dimitris Papailiopoulos, and Kangwook Lee. 2024. [Predictive Pipelined Decoding: A Compute-Latency Trade-off for Exact LLM Decoding](#). *Transactions on Machine Learning Research*.

Jun Zhang, Jue Wang, Huan Li, Lidan Shou, Ke Chen, Gang Chen, and Sharad Mehrotra. 2024a. [Draft&Verify: Lossless Large Language Model Acceleration via Self-Speculative Decoding](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11263–11282, Bangkok, Thailand. Association for Computational Linguistics.

Songming Zhang, Xue Zhang, Zengkui Sun, Yufeng Chen, and Jinan Xu. 2024b. [Dual-space knowledge distillation for large language models](#). In *Proceedings of the 2024 Conference on Empirical Methods in*

*Natural Language Processing*, pages 18164–18181, Miami, Florida, USA. Association for Computational Linguistics.

Yilong Zhao, Chien-Yu Lin, Kan Zhu, Zihao Ye, Lequn Chen, Size Zheng, Luis Ceze, Arvind Krishnamurthy, Tianqi Chen, and Baris Kasikci. 2024. Atom: Low-bit quantization for efficient and accurate llm serving. *Proceedings of Machine Learning and Systems*, 6:196–209.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, and others. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36:46595–46623.

Yongchao Zhou, Kaifeng Lyu, Ankit Singh Rawat, Aditya Krishna Menon, Afshin Rostamizadeh, Sanjiv Kumar, Jean-François Kagy, and Rishabh Agarwal. 2024. [DistillSpec: Improving Speculative Decoding via Knowledge Distillation](#). In *The Twelfth International Conference on Learning Representations*.

## A More on Speculative Decoding

Autoregressive decoding (Touvron et al., 2023; Bai et al., 2023; Jiang et al., 2023; OpenAI et al., 2024), suffers from inefficiency because it generates text one token at a time (Figure 5, Left). Speculative decoding (Chen et al., 2023; Leviathan et al., 2023) follows a *guess-and-verify* paradigm (Figure 5, Right). In speculative decoding, a smaller LLM (draft model) (Chen et al., 2023; Leviathan et al., 2023; Miao et al., 2024; Sun et al., 2023b; Zhou et al., 2024; Cai et al., 2024) or the original LLM trained in a specialized manner (self-speculative decoding) (Elhoushi et al., 2024; Liu et al., 2024a; Yang et al., 2024; Zhang et al., 2024a; Li et al., 2024b) predicts multiple tokens in advance. The original LLM then verifies these predictions in parallel, improving efficiency.

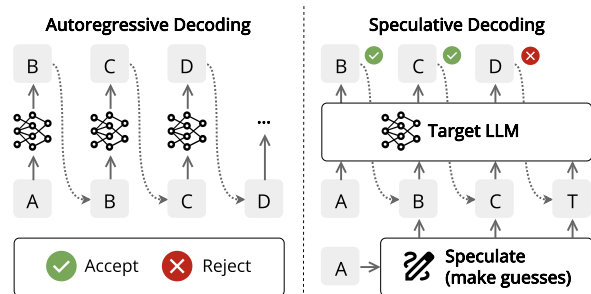


Figure 5: Examples of Autoregressive decoding (Left) and Speculative Decoding (Right). While autoregressive decoding generates one token per forward step, speculative decoding generates three tokens with one forward step.



LLMs process discrete integer sequences as inputs, where each integer represents a token. We define the input sequence as  $\mathbf{x} = (x_1, x_2, \dots, x_s) \in \mathbb{N}^s$  of length  $s$ , and denote a slice of length  $m$  at step  $t$  as  $\mathbf{x}_{1:m} = (x_1, x_2, \dots, x_m)$ . The output of an LLM represents the probability distribution over the next token. The probability of generating the  $s$ -th token, conditioned on all preceding tokens, is given by  $P_M(x_s \mid x_{1:s-1})$ . The next token  $x_s$  is then sampled from this distribution using various methods (e.g., greedy, top- $k$ , and top- $p$  sampling; see (Kool et al., 2020; Holtzman et al., 2020)). In the case of greedy sampling, the next token is selected as  $x_s = \arg \max P_M(x_s \mid x_{1:s-1})$ .

Let  $\mathbf{x}^0$  be the prompt tokens provided by the user. The LLM generates an output sequence of length  $m$ , with each generated token  $y_i$  computed autoregressively. Assuming greedy sampling, the decoding process follows:

$$\begin{cases} y_1 = \arg \max P_M(y_1 \mid \mathbf{x}) \\ y_2 = \arg \max P_M(y_2 \mid y_1, \mathbf{x}) \\ \vdots \\ y_m = \arg \max P_M(y_m \mid y_{1:m-1}, \mathbf{x}). \end{cases} \quad (1)$$

### A.1 Speculative Decoding

Speculative decoding follows a *Guess-And-Verify* approach, where multiple candidate future tokens are speculated and subsequently verified in a single decoding step. With tree attention (Miao et al., 2024), multiple drafts can be verified simultaneously. Let  $G$  denote the number of guesses, and define the set of guesses as  $\tilde{Y} = \{\tilde{y}^{(0)}, \tilde{y}^{(1)}, \dots, \tilde{y}^{(G)}\}$ , where each guess sequence has length  $K$ . The  $j$ -th token of the  $i$ -th guess is denoted as  $\tilde{y}_j^{(i)}$ .

In the case of speculative decoding with greedy sampling, given the prompt  $\mathbf{x}$ , a drafting method is used to generate the draft sequences  $\tilde{Y}$ . Using these drafts, the LLM then computes the true tokens  $(y'_1, y'_2, \dots, y'_K)$  in parallel. For instance, for the guess sequence  $\tilde{y}^{(0)}$ , the true tokens are determined as:

$$\begin{cases} y'_1 = \arg \max P_M(y_1 \mid \mathbf{x}) \\ y'_2 = \arg \max P_M(y_2 \mid \tilde{y}_1^{(0)}, \mathbf{x}) \\ \vdots \\ y'_K = \arg \max P_M(y_K \mid \tilde{y}_{1:K-1}^{(0)}, \mathbf{x}). \end{cases} \quad (2)$$

These generated tokens are then verified. Let  $h$  be the highest number of correct guessed tokens

across all guesses. Consequently,  $h + 1$  tokens are generated in one forward step. Algorithm 2 outlines speculative decoding with greedy sampling.

## B Implementation Details

### B.1 Frameworks and Libraries

We implement SPECTRA in Python using PyTorch 2.1.0 and the Hugging Face transformers library. For large-scale model loading (e.g., LLaMA-2-70B, LLaMA-3-70B), we employ 16-bit (FP16) precision with a pre-allocated key-value cache.

### B.2 Models and Checkpoints

We run our experiments primarily with:

- **LLaMA-2-Chat** (Touvron et al., 2023) in sizes 7B, 13B, 70B.
- **CodeLlama** (Rozière et al., 2024) in sizes 7B and 13B.
- **LLaMA-3-Instruct** (Dubey et al., 2024) in sizes 8B and 70B.

All checkpoints are obtained from the official or Hugging Face repositories without fine-tuning or modification. For each model, we enable half-precision inference. We also verify numerically (by comparing 32-bit and 16-bit outputs) that speculative decoding preserves exact or near-exact token sequences within typical floating-point tolerances.

### B.3 Hardware

Most experiments are conducted on a single NVIDIA A100 GPU with 80GB of memory. We also evaluate on other NVIDIA GPUs (RTX 3090, RTX 8000, A40, A6000) to study hardware-specific scaling. For the largest checkpoints (70B) that do not fit on a single GPU under certain configurations, we optionally distribute them across multiple GPUs (2x, 4x, or 8x H100) using standard pipeline-parallelism from Hugging Face’s library.

### B.4 Hyperparameters

**Lookahead, REST, and ANPD.** We replicate each baseline using their publicly available GitHub code, keeping to the default settings and hyperparameters outlined in the original papers.

**Spectra.** By default, we use a 5-gram setup for our forward/backward dictionaries, storing all subsequences (i.e., sub-ngrams). We also maintain a *candidate pool* of size  $\mathbf{W} = 15$  per key to generate

new n-gram records; after each forward pass, candidate sequences are shifted by one token and then re-populated. We introduce a threshold  $\tau \in [0, 1]$ , default set to 0.1, to decide when to force the selection of a token not yet in the forward dictionary. This mechanism balances coverage of unseen prefixes with reinforcing common contexts. At every speculative decoding step, we allow up to  $G = 60$  guess tokens. Internal guesses receive priority, and if there is still capacity under the guess limit, we add external guesses.

For external lookups, we implement a trie structure for rapid prefix queries, following a design similar to REST (He et al., 2024). For **conversation** tasks (e.g., MT-Bench), we gather approximately 100k examples from the UltraChat dataset (Ding et al., 2023), focusing on those with minimal perplexity under the *same* LLM we aim to accelerate. For **code** tasks (e.g., HumanEval, MBPP), we draw from TheStack (Kocetkov et al., 2023) and again refine it to the 100k snippets with the lowest perplexity for memory efficiency. We measure perplexity by running a single forward pass (in streaming mode) over candidate samples and ranking them. Despite being relatively small (100k), this curated corpus achieves robust guess quality.

All speedup and throughput metrics are computed at a batch size of 1. In code generation tasks, the maximum generation length is typically 512 tokens, whereas for conversation tasks (MT-Bench, GSM8K), we allow up to 1024 tokens or stop early if the model outputs an end-of-sequence token. All random seeds are set to 0.

## C Details Results with Throughputs

We provide a detailed throughput analysis to complement the speedup ratios reported in the main text. Our goal is to demonstrate how SPECTRA scales across various model sizes, datasets, and GPU architectures. We measure throughput using two key metrics:

- **Macro Throughput (Mac-TP).** Calculated as the average of per-generation token-processing rates—i.e., for each generation step  $i$ , we compute  $token_i/time_i$  and then average over all steps.
- **Micro Throughput (Mic-TP).** Calculated as the total number of generated tokens divided by the total elapsed time

Table 4 focuses on GSM8K and MTBench performance across four different GPU models, while Table 3 provides more granular results on additional datasets and model configurations. In all cases, SPECTRA consistently achieves higher throughput than both non-speculative baselines and other training-free accelerators, as evidenced by improvements in both Mic-TP and Mac-TP. Notably, this performance advantage remains stable even on older GPUs (e.g., the RTX 3090 and RTX 8000), demonstrating SPECTRA’s robustness to varying hardware capabilities.

## D Evaluating SPECTRA in Different GPU Types

**Different GPU types.** Table 5 reports speedups on GSM8K and MTBench across four GPUs with varying memory throughput and compute capabilities. While absolute wall-clock times differ across GPUs, the *relative* accelerations remain consistent. SPECTRA consistently outperforms other baselines, including Lookahead, achieving higher speedups in all cases. On older GPUs (e.g., RTX 3090 or RTX 8000), the gap between Lookahead and SPECTRA narrows slightly due to less efficient parallelism, but SPECTRA maintains its lead. These results demonstrate that SPECTRA is robust to hardware variations and effective across both data-center and consumer-grade GPUs.

## E Evaluating SPECTRA in Multi-GPU Environments

A critical consideration for practical deployment is how SPECTRA scales when models are distributed across multiple GPUs—a common requirement for large LLMs exceeding single-device memory capacity. To evaluate this, we measure SPECTRA’s performance under three distributed configurations of LLaMA-2-70B: (1) 2xH100 with full precision, (2) 4xH100 with full precision, and (3) 8xH100 with full precision. We also include a baseline of 1xH100 with 8-bit quantization for memory-constrained single-GPU inference. Table 6 reports throughput and speedup metrics.

SPECTRA achieves consistent speedups of 2.00— $2.03\times$  across all multi-GPU configurations while maintaining a stable compression ratio ( $\tau$ ) of 2.52. This demonstrates robust scalability—partitioning model weights introduces minimal overhead, and the speculative verification process remains efficient despite inter-GPU communication. Notably,

Model	Method	Classeval		GSM8K		Humaneval		MBPP		MTBench	
		Mac-TP	Mic-TP	Mac-TP	Mic-TP	Mac-TP	Mic-TP	Mac-TP	Mic-TP	Mac-TP	Mic-TP
Greedy (temperature=0)											
CL-13B	Autoregressive	30.85	30.85	32.03	32.03	32.35	32.35	32.07	32.07	30.69	30.63
CL-13B	ANPD	59.77	58.03	89.99	89.18	67.43	64.65	86.76	86.41	80.10	76.68
CL-13B	Lookahead	69.28	68.62	89.73	89.00	74.33	73.23	93.38	92.80	79.38	78.67
CL-13B	REST	39.53	37.73	29.93	29.47	51.15	47.49	27.41	27.39	28.92	27.18
CL-13B	SPECTRA (Ours)	<b>73.47</b>	<b>72.98</b>	<b>93.36</b>	<b>93.23</b>	<b>84.91</b>	<b>84.41</b>	<b>105.44</b>	<b>105.39</b>	<b>81.32</b>	<b>80.68</b>
CL-7B	Autoregressive	41.17	41.17	41.17	41.17	41.41	41.41	41.60	41.60	38.91	38.93
CL-7B	ANPD	94.76	93.02	132.26	131.30	89.26	87.13	131.35	130.99	130.41	126.64
CL-7B	Lookahead	106.51	105.95	123.04	121.90	103.45	103.51	120.75	120.23	125.58	124.77
CL-7B	REST	59.49	56.61	37.61	37.21	70.38	65.22	40.11	40.09	39.64	36.70
CL-7B	SPECTRA (Ours)	<b>111.09</b>	<b>110.68</b>	<b>137.24</b>	<b>136.86</b>	<b>122.54</b>	<b>122.41</b>	<b>148.32</b>	<b>148.07</b>	<b>143.98</b>	<b>144.32</b>
L2-13B	Autoregressive	31.85	31.56	32.40	32.43	32.27	32.27	32.19	32.19	31.93	31.78
L2-13B	ANPD	43.30	44.44	47.54	45.22	43.24	42.28	36.20	35.84	37.44	34.84
L2-13B	Lookahead	57.49	58.94	47.44	47.62	55.76	55.58	44.41	44.15	48.11	46.62
L2-13B	REST	38.81	37.74	30.36	30.22	40.47	39.70	30.70	30.67	36.39	37.02
L2-13B	SPECTRA (Ours)	<b>63.64</b>	<b>64.31</b>	<b>59.21</b>	<b>58.63</b>	<b>63.39</b>	<b>63.18</b>	<b>52.43</b>	<b>52.19</b>	<b>56.04</b>	<b>53.75</b>
L2-70B	Autoregressive	2.60	2.60	2.61	2.61	2.61	2.61	2.63	2.63	2.60	2.60
L2-70B	ANPD	4.72	4.80	4.25	4.10	4.85	4.76	3.07	3.07	3.47	3.30
L2-70B	Lookahead	6.90	7.16	4.87	5.12	6.71	6.73	3.92	3.93	5.05	5.02
L2-70B	SPECTRA (Ours)	<b>8.07</b>	<b>8.35</b>	<b>6.58</b>	<b>6.75</b>	<b>8.41</b>	<b>8.41</b>	<b>4.88</b>	<b>4.88</b>	<b>6.32</b>	<b>6.22</b>
L2-7B	Autoregressive	40.33	40.32	41.01	41.03	41.14	41.13	41.00	41.04	40.48	40.50
L2-7B	ANPD	65.54	68.10	62.40	59.38	63.27	59.98	48.94	47.67	52.47	50.06
L2-7B	Lookahead	88.41	91.05	68.00	68.20	84.69	83.87	59.79	60.76	70.04	69.07
L2-7B	REST	54.74	53.93	41.43	41.38	57.99	56.41	41.28	40.74	50.58	51.79
L2-7B	SPECTRA (Ours)	<b>96.88</b>	<b>98.75</b>	<b>86.51</b>	<b>85.50</b>	<b>98.77</b>	<b>98.38</b>	<b>72.39</b>	<b>73.22</b>	<b>81.93</b>	<b>79.20</b>
L3-70B	Autoregressive	2.58	2.57	2.58	2.58	2.59	2.59	2.59	2.59	2.55	2.55
L3-70B	ANPD	3.97	4.19	3.86	3.72	4.72	4.75	3.77	3.59	3.14	3.03
L3-70B	Lookahead	6.17	6.47	3.99	3.96	6.63	6.75	3.70	3.66	4.49	4.53
L3-70B	SPECTRA (Ours)	<b>6.87</b>	<b>7.18</b>	<b>5.43</b>	<b>5.34</b>	<b>7.33</b>	<b>7.50</b>	<b>5.01</b>	<b>4.88</b>	<b>5.25</b>	<b>5.16</b>
L3-8B	Autoregressive	36.59	36.58	36.74	36.74	36.20	36.21	35.24	35.20	36.55	36.69
L3-8B	ANPD	77.21	78.76	141.89	141.36	66.31	65.57	118.47	112.95	41.77	40.20
L3-8B	Lookahead	94.92	97.09	136.32	135.92	89.99	90.47	133.67	133.12	56.09	55.49
L3-8B	SPECTRA (Ours)	<b>103.61</b>	<b>105.88</b>	<b>142.89</b>	<b>142.72</b>	<b>92.86</b>	<b>93.16</b>	<b>143.80</b>	<b>142.72</b>	<b>61.69</b>	<b>60.22</b>
Sampling (temperature=1.0)											
CL-13B	Autoregressive	30.90	30.64	31.38	31.37	31.24	31.39	31.46	31.45	30.71	30.67
CL-13B	ANPD	35.48	34.86	33.54	32.34	32.64	34.36	31.57	30.95	70.92	65.68
CL-13B	Lookahead	42.54	40.74	33.79	32.49	40.25	42.17	32.02	31.19	71.50	68.46
CL-13B	REST	35.15	33.22	25.67	25.24	39.58	38.49	26.43	25.89	28.41	26.69
CL-13B	SPECTRA (Ours)	<b>51.86</b>	<b>50.04</b>	<b>37.57</b>	<b>35.67</b>	<b>51.60</b>	<b>52.64</b>	<b>36.29</b>	<b>35.27</b>	<b>72.90</b>	<b>69.98</b>
CL-7B	Autoregressive	39.60	39.58	40.85	40.87	40.05	40.10	40.81	40.81	40.49	40.50
CL-7B	ANPD	50.89	51.76	47.44	46.68	44.14	46.34	45.86	45.81	112.29	103.57
CL-7B	Lookahead	60.87	60.29	48.54	47.64	57.12	61.14	48.64	48.27	110.07	105.00
CL-7B	REST	48.64	46.41	35.98	35.46	53.35	52.26	37.04	36.57	39.36	36.51
CL-7B	SPECTRA (Ours)	<b>71.70</b>	<b>71.78</b>	<b>55.24</b>	<b>52.81</b>	<b>67.27</b>	<b>69.20</b>	<b>54.48</b>	<b>52.91</b>	<b>112.43</b>	<b>108.49</b>
L2-13B	Autoregressive	31.23	31.17	31.44	31.47	31.41	31.42	32.02	32.06	31.67	31.59
L2-13B	ANPD	37.53	37.94	39.11	37.99	36.79	36.75	32.97	32.71	36.91	34.34
L2-13B	Lookahead	47.59	47.35	41.60	41.76	46.33	46.51	37.82	37.82	47.35	45.48
L2-13B	REST	36.78	36.17	29.33	29.25	37.46	36.71	29.38	29.28	35.50	36.21
L2-13B	SPECTRA (Ours)	<b>53.13</b>	<b>52.28</b>	<b>48.60</b>	<b>48.11</b>	<b>52.93</b>	<b>53.11</b>	<b>42.95</b>	<b>43.03</b>	<b>54.98</b>	<b>52.42</b>
L2-7B	Autoregressive	39.89	39.88	40.58	40.59	40.09	40.10	40.59	40.66	40.65	40.70
L2-7B	ANPD	52.14	52.78	54.23	52.90	51.40	50.97	44.73	43.77	50.92	48.24
L2-7B	Lookahead	70.82	71.17	61.15	61.34	68.78	69.01	50.84	51.83	68.27	66.77
L2-7B	REST	50.35	49.99	40.19	40.09	50.86	50.06	38.94	38.18	49.12	50.54
L2-7B	SPECTRA (Ours)	<b>78.46</b>	<b>78.74</b>	<b>72.13</b>	<b>71.68</b>	<b>81.71</b>	<b>81.76</b>	<b>59.77</b>	<b>60.09</b>	<b>80.21</b>	<b>77.00</b>
L3-8B	Autoregressive	35.75	35.76	35.16	35.17	36.01	36.02	36.05	36.07	35.39	35.48
L3-8B	ANPD	44.71	43.72	69.12	66.73	51.48	51.57	68.03	64.54	40.84	39.23
L3-8B	Lookahead	53.05	50.57	72.68	69.11	64.59	63.79	71.88	68.90	55.46	53.74
L3-8B	SPECTRA (Ours)	<b>69.50</b>	<b>68.92</b>	<b>79.88</b>	<b>76.53</b>	<b>69.09</b>	<b>68.62</b>	<b>78.99</b>	<b>76.69</b>	<b>60.33</b>	<b>57.69</b>

Table 3: Micro throughput (Mic-TP) and Macro throughput (Mac-TP) across multiple tasks and models.

GPU	Method	GSM8K		MTBench	
		Mac-TP	Mic-TP	Mac-TP	Mic-TP
A40	Autoregressive	32.66	32.66	32.14	31.66
	Lookahead	48.59	48.73	49.13	47.96
	SPECTRA	<b>62.56</b>	<b>61.52</b>	<b>59.00</b>	<b>56.80</b>
A6000	Autoregressive	39.15	39.17	38.78	38.24
	Lookahead	58.13	58.30	58.84	57.40
	SPECTRA	<b>75.20</b>	<b>74.16</b>	<b>71.3</b>	<b>69.28</b>
RTX8000	Autoregressive	34.03	34.27	34.21	34.02
	Lookahead	45.25	45.42	45.73	44.16
	SPECTRA	<b>57.95</b>	<b>57.09</b>	<b>54.16</b>	<b>52.32</b>
RTX3090	Autoregressive	40.67	40.76	41.17	41.22
	Lookahead	53.69	53.75	53.51	52.09
	SPECTRA	<b>74.87</b>	<b>73.88</b>	<b>71.58</b>	<b>69.79</b>

Table 4: Throughput results for different GPU types on GSM8K and MTBench.

GPU	Method	GSM8K		MTBench	
		speedup	$\tau$	speedup	$\tau$
A40	Lookahead	1.49	1.93	1.53	2.07
	SPECTRA	<b>1.92</b>	<b>2.46</b>	<b>1.84</b>	<b>2.36</b>
A6000	Lookahead	1.48	1.92	1.52	2.06
	SPECTRA	<b>1.92</b>	<b>2.46</b>	<b>1.84</b>	<b>2.36</b>
RTX8000	Lookahead	1.33	1.93	1.34	2.08
	SPECTRA	<b>1.70</b>	<b>2.46</b>	<b>1.58</b>	<b>2.35</b>
RTX3090	Lookahead	1.32	1.92	1.30	2.06
	SPECTRA	<b>1.84</b>	<b>2.46</b>	<b>1.74</b>	<b>2.36</b>

Table 5: Hardware scalability of SPECTRA decoding on GSM8K and MTBench for various GPU architectures.

even in the quantized single-GPU setting, SPECTRA provides a  $2.43\times$  speedup, outperforming standard autoregressive decoding. These results validate SPECTRA’s practicality for large-scale deployments where memory constraints necessitate distributed inference.

## F Verifying Generation Quality with SPECTRA Decoding

**Greedy Decoding Performance.** To assess the quality of greedy decoding, we compare the inference results of the LLaMA-2-7B Chat model using SPECTRA Decoding against Hugging Face’s standard greedy search. Our baseline consists of single-precision (FP32) inference on 160 conversational turns from the MT-Bench dataset. Under FP32, SPECTRA Decoding produces identical outputs to the baseline.

However, when transitioning to half-precision (FP16), even Hugging Face’s native greedy search

generates 25 discrepancies (out of 160) compared to the FP32 baseline. SPECTRA Decoding exhibits a similar discrepancy rate (26), confirming that it maintains the output distribution within the numerical error margins typically observed in standard half-precision inference libraries.

**Sampling Decoding Performance.** We also assess generation quality under a stochastic sampling setting (temperature = 1.0). As detailed in Table 7, SPECTRA Decoding produces ROUGE-1, ROUGE-2, and ROUGE-L scores on both the CNN/DailyMail (Nallapati et al., 2016) and XSum (Narayan et al., 2018) summarization datasets that are nearly identical to those of standard autoregressive sampling. At the same time, SPECTRA achieves notable speedups ( $1.60\times$  on CNN/DailyMail and  $1.69\times$  on XSum) with compression ratios of 2.05 and 2.08, respectively. These results confirm that SPECTRA Decoding accelerates inference while preserving generation quality across diverse tasks.

These findings reaffirm that SPECTRA Decoding, does not degrade generation quality compared to conventional greedy or sampling-based methods.

## G Token Acceptance Rate Analysis

Figure 6 plots the cumulative number of accepted tokens versus decoding steps for each dataset (MTBench, HumanEval, MBPP, and GSM8K). The steeper ascent of the SPECTRA curve indicates that our method requires substantially fewer decoding steps compared to alternatives, for example, almost two times shorter than ANPD. This improvement is



GPU & Model Setting	Method	MTBench			
		Mac-TP	Mic-TP	Speedup	$\tau$
1xH100 - Quantized Int8	Autoregressive	2.60	2.60	1.00	1.00
	SPECTRA	<b>6.32</b>	<b>6.22</b>	<b>2.43</b>	<b>2.51</b>
2xH100 - FP16	Autoregressive	14.81	14.70	1.00	1.00
	SPECTRA	<b>29.62</b>	<b>28.91</b>	<b>2.00</b>	<b>2.52</b>
4xH100 - FP16	Autoregressive	14.60	14.48	1.00	1.00
	SPECTRA	<b>29.67</b>	<b>28.89</b>	<b>2.03</b>	<b>2.52</b>
8xH100 - FP16	Autoregressive	14.39	14.28	1.00	1.00
	SPECTRA	<b>29.27</b>	<b>28.55</b>	<b>2.03</b>	<b>2.52</b>

Table 6: Results in multi-GPU Enviroments on GSM8K and MTBench using LLama-2-chat-70B.

Dataset	Method	ROUGE-1	ROUGE-2	ROUGE-L	Speedup	$\tau$
CNN	Autoregressive	9.77	0.39	7.20	1.00	1.00
	SPECTRA	9.74	0.41	7.18	1.60	2.05
XSUM	Autoregressive	18.12	4.36	12.43	1.00	1.00
	SPECTRA	18.13	4.40	12.49	1.69	2.08

Table 7: Evaluation of SPECTRA Decoding on CNN/DailyMail and XSum using a temperature of 1.0. ROUGE scores, speedups over autoregressive decoding, and compression ratio ( $\tau$ ) are reported for LLaMA-2-7B-Chat.

attributed to a higher token acceptance rate, which in turn reduces the overall number of decoding iterations and enhances the efficiency of the generation process.

## H Algorithms

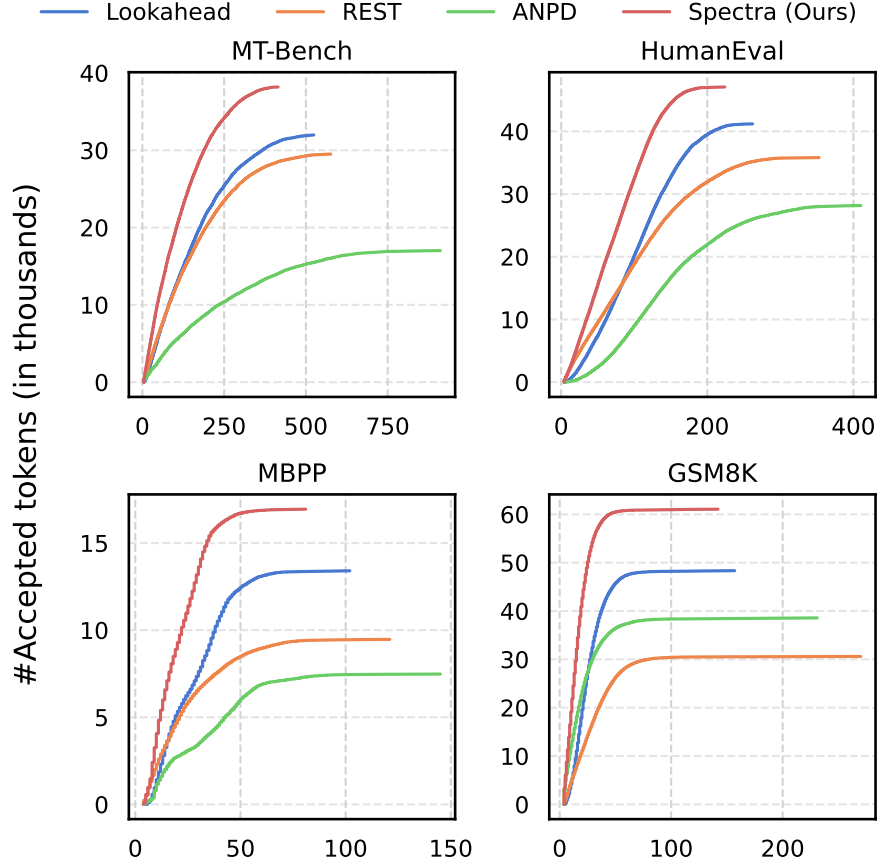


Figure 6: Total number of accepted tokens across all samples at each decoding step.

---

**Algorithm 2** Speculative Decoding (Multiple guesses and Greedy Sampling)

---

**Given** guess size  $K$ , number of guesses  $G$ , and target length  $T$ .

**Given** initial prompt sequence  $\mathbf{x}$ .

**while**  $n < T$  **do**

    Obtain multiple drafts  $\tilde{Y} = \{\tilde{y}^{(0)}, \tilde{y}^{(1)}, \dots, \tilde{y}^{(G)}\}$ .

    In parallel, compute  $K + 1$  verification tokens  $y'$ :

**for**  $i = 1 : K$  **do**

$y_i^{(g)} = \arg \max P_M(y_i | \tilde{y}_{i-1}^{(g)}, \mathbf{x}), \quad \forall g \in \{0, \dots, G\}$

**end for**

    Identify the sequence  $\tilde{y}^{(g^*)}$  with the highest token matches and the corresponding  $y'^{(g^*)}$ .

**for**  $t = 1 : K$  **do**

**if**  $y_t^{(g)} = \tilde{y}_t^{(g^*)}$  **then**

            Set  $y_{n+t} \leftarrow \tilde{y}_t^{(g^*)}$  and  $n \leftarrow n + 1$ .

**else**

$y_{n+t} \leftarrow y_t^{(g)}$  and exit for loop.

**end if**

**end for**

**end while**

---

---

**Algorithm 3** Greedy Verification with SPECTRA DECODING

---

**Require:** sequence  $x$ , model  $P_M$ , guesses  $g^i$  with  $i \in [1, G]$

**Ensure:**  $o$  {accepted tokens of length 1 to  $N$ }

```
1: function GREEDYVERIFICATION( $x, P_M, g$ )
2:    $V, D, o \leftarrow \emptyset, \emptyset, \emptyset$ 
3:   for  $i = 1$  to  $G$  do
4:      $V.append(g_2^i)$  ▷ each  $g_2^i$  is a n-1 gram
5:      $D.append(P_M(g_2^i, x_{\text{next}}|g_2^i, x))$  ▷ obtain last token of  $x$  and all  $g_2^i$ 's outputs – totally  $N$  distributions
6:   end for
7:   for  $i = 1$  to  $N - 1$  do
8:      $j \leftarrow 1$ 
9:      $\text{is\_accept} \leftarrow 0$ 
10:     $\mathcal{P} \leftarrow D[l][i]$ 
11:    while  $j \leq \text{size}(V)$  do
12:       $s_j \leftarrow V[j]$ 
13:      if  $s_j = \arg \max \mathcal{P}$  then ▷ accepted, update all potential speculations and probabilities
14:         $o.append(s_j)$ 
15:         $\text{is\_accept} \leftarrow 1$ 
16:         $V_{\text{new}}, D_{\text{new}} \leftarrow \emptyset, \emptyset$ 
17:        for  $k = j$  to  $\text{size}(V)$  do
18:          if  $s_j = V[k]$  then
19:             $V_{\text{new}}.append(V[k])$ 
20:             $D_{\text{new}}.append(D[k])$ 
21:          end if
22:        end for
23:         $V, D \leftarrow V_{\text{new}}, D_{\text{new}}$ 
24:        break
25:      else ▷ rejected, go to next speculation
26:         $j \leftarrow j + 1$ 
27:      end if
28:    end while
29:    if  $\text{is\_accept}$  then
30:      continue
31:    else ▷ guarantee one step movement
32:       $o.append(\arg \max \mathcal{P})$ 
33:      break
34:    end if
35:  end for
36:  if  $\text{is\_accept}$  then
37:     $o.append(\arg \max D[1]_N)$ 
38:  end if
39:  return  $o$ 
40: end function
```

---

---

**Algorithm 4** Sample Verification with SPECTRA DECODING

---

**Require:** sequence  $x$ , model  $P_M$ , guesses  $g^i$  with  $i \in [1, G]$

**Ensure:**  $o$  {accepted tokens of length 1 to  $N$ }

```
1: function SAMPLEVERIFICATION( $x, P_M, g$ )
2:    $V, D, o \leftarrow \emptyset, \emptyset, \emptyset$ 
3:   for  $i = 1$  to  $G$  do
4:      $V.append(g_2^i)$  ▷ each  $g_2^i$  is a n-1 gram
5:      $D.append(P_M(g_2^i, x_{\text{next}}|g_2^i, x))$  ▷ obtain last token of  $x^0$  and all  $g_2^i$ 's outputs – totally  $N$  probability distributions
6:   end for
7:   for  $i = 1$  to  $N - 1$  do
8:      $j \leftarrow 1$ 
9:      $\text{is\_accept} \leftarrow 0$ 
10:     $\mathcal{P}_j \leftarrow D[1]_i$ 
11:    while  $j \leq \text{size}(V)$  do
12:       $s_j \leftarrow V[j]$ 
13:      sample  $r \sim U(0, 1)$ 
14:      if  $r \leq \mathcal{P}_j(s_j)$  then ▷ accepted, update all potential speculations and probabilities
15:         $o.append(s_j)$ 
16:         $\text{is\_accept} \leftarrow 1$ 
17:         $V_{\text{new}}, D_{\text{new}} \leftarrow \emptyset, \emptyset$ 
18:        for  $k = j$  to  $\text{size}(V)$  do
19:          if  $s_j = V[k]$  then
20:             $V_{\text{new}}.append(V[k])$ 
21:             $D_{\text{new}}.append(D[k])$ 
22:          end if
23:        end for
24:         $V, D \leftarrow V_{\text{new}}, D_{\text{new}}$ 
25:        break
26:      else ▷ rejected, go to next speculation
27:         $\mathcal{P}_j(s_j) \leftarrow 0$ 
28:         $\mathcal{P}_{j+1} = \text{norm}(\mathcal{P}_j)$ 
29:         $j \leftarrow j + 1$ 
30:      end if
31:    end while
32:    if  $\text{is\_accept}$  then
33:      continue
34:    else ▷ guarantee one step movement
35:      sample  $x_{\text{next}} \sim \mathcal{P}_j$ 
36:       $o.append(x_{\text{next}})$ 
37:      break
38:    end if
39:  end for
40:  if  $\text{is\_accept}$  then
41:     $o.append(\text{sample } x_{\text{next}} \sim D[1]_N)$ 
42:  end if
43:  return  $o$ 
44: end function
```

---